




COMP20003
Algorithms and Data Structures

Shortest Paths

Nir Lipovetzky
Department of Computing and Information Systems
University of Melbourne
Semester 2



https://commons.wikimedia.org/wiki/File:Haas_Hel_Hexacoralla.jpg



Example weighted graph


Adjacency List

0 → 1 → 2

1 → 0 → 2 → 3

2 → 0 → 1 → 3

3 → 1 → 2




Previous visit order from node 0:

- But if these are restaurants and bars, and we want to go to a **nearby** bar From restaurant 0...
... in this case the answer is easy. But if you scale it...




Example weighted graph



No direct route, several possible paths



Example weighted graph bfs



Length of a path:

Sum of edges along the way.

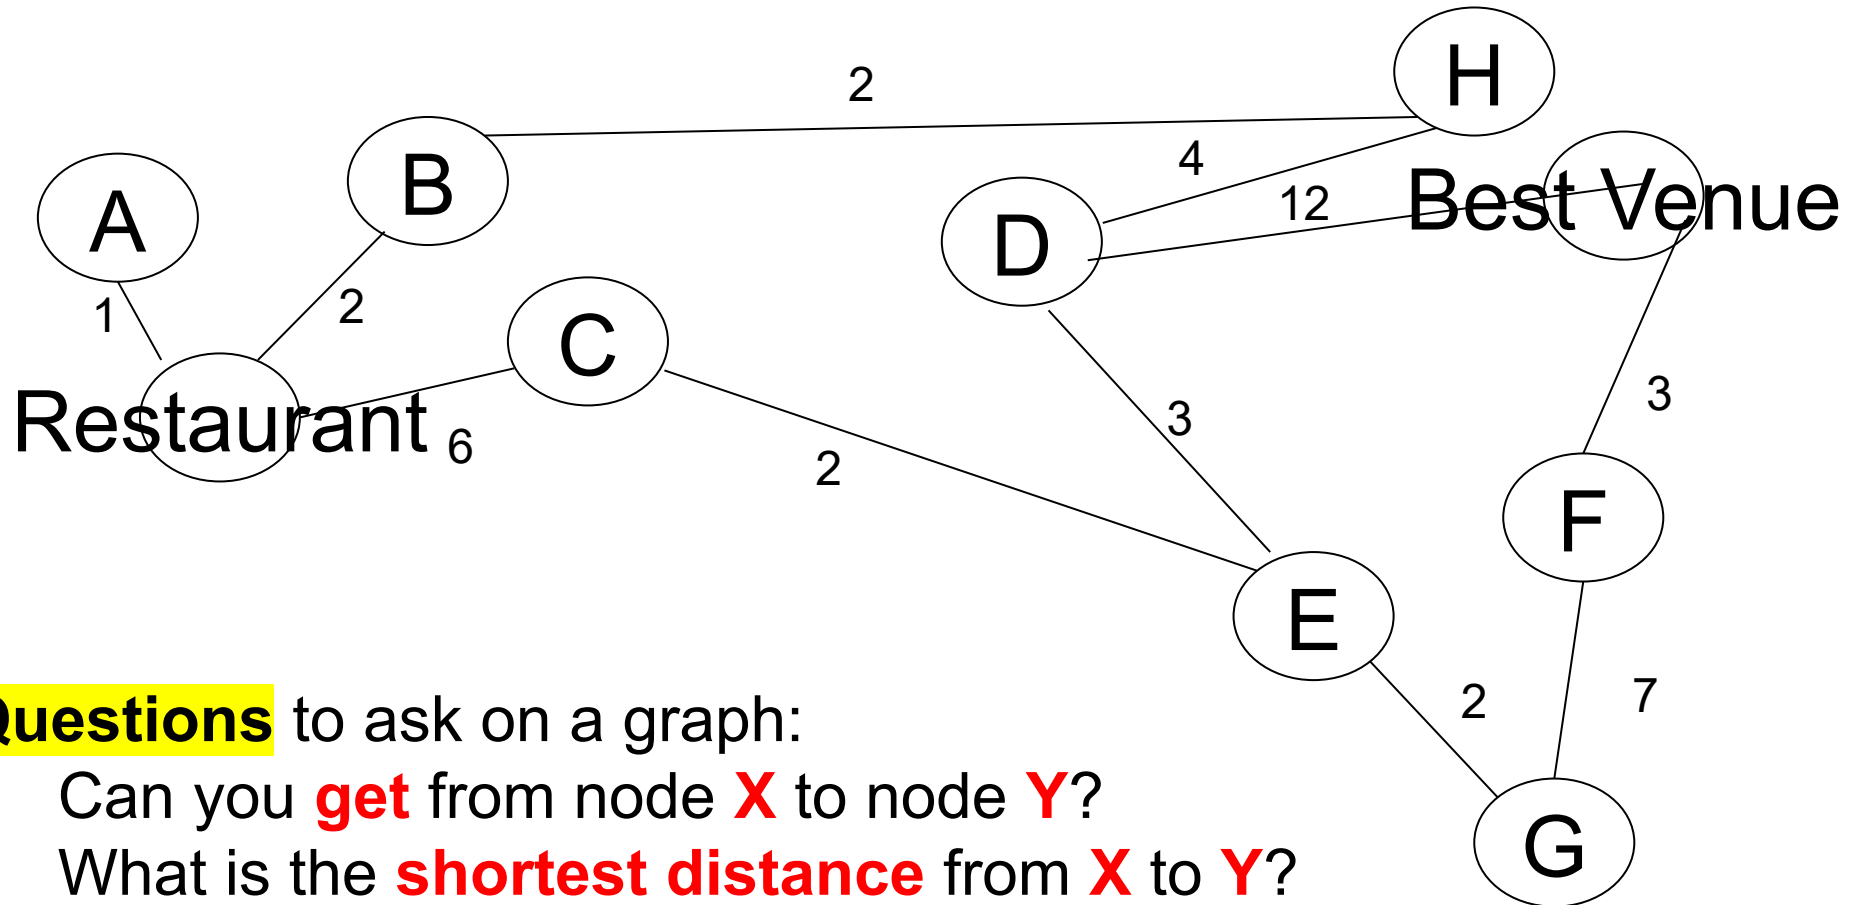
Length of R->B->H->D->Best

Distance between nodes:

Length of **shortest path**.



Example weighted graph

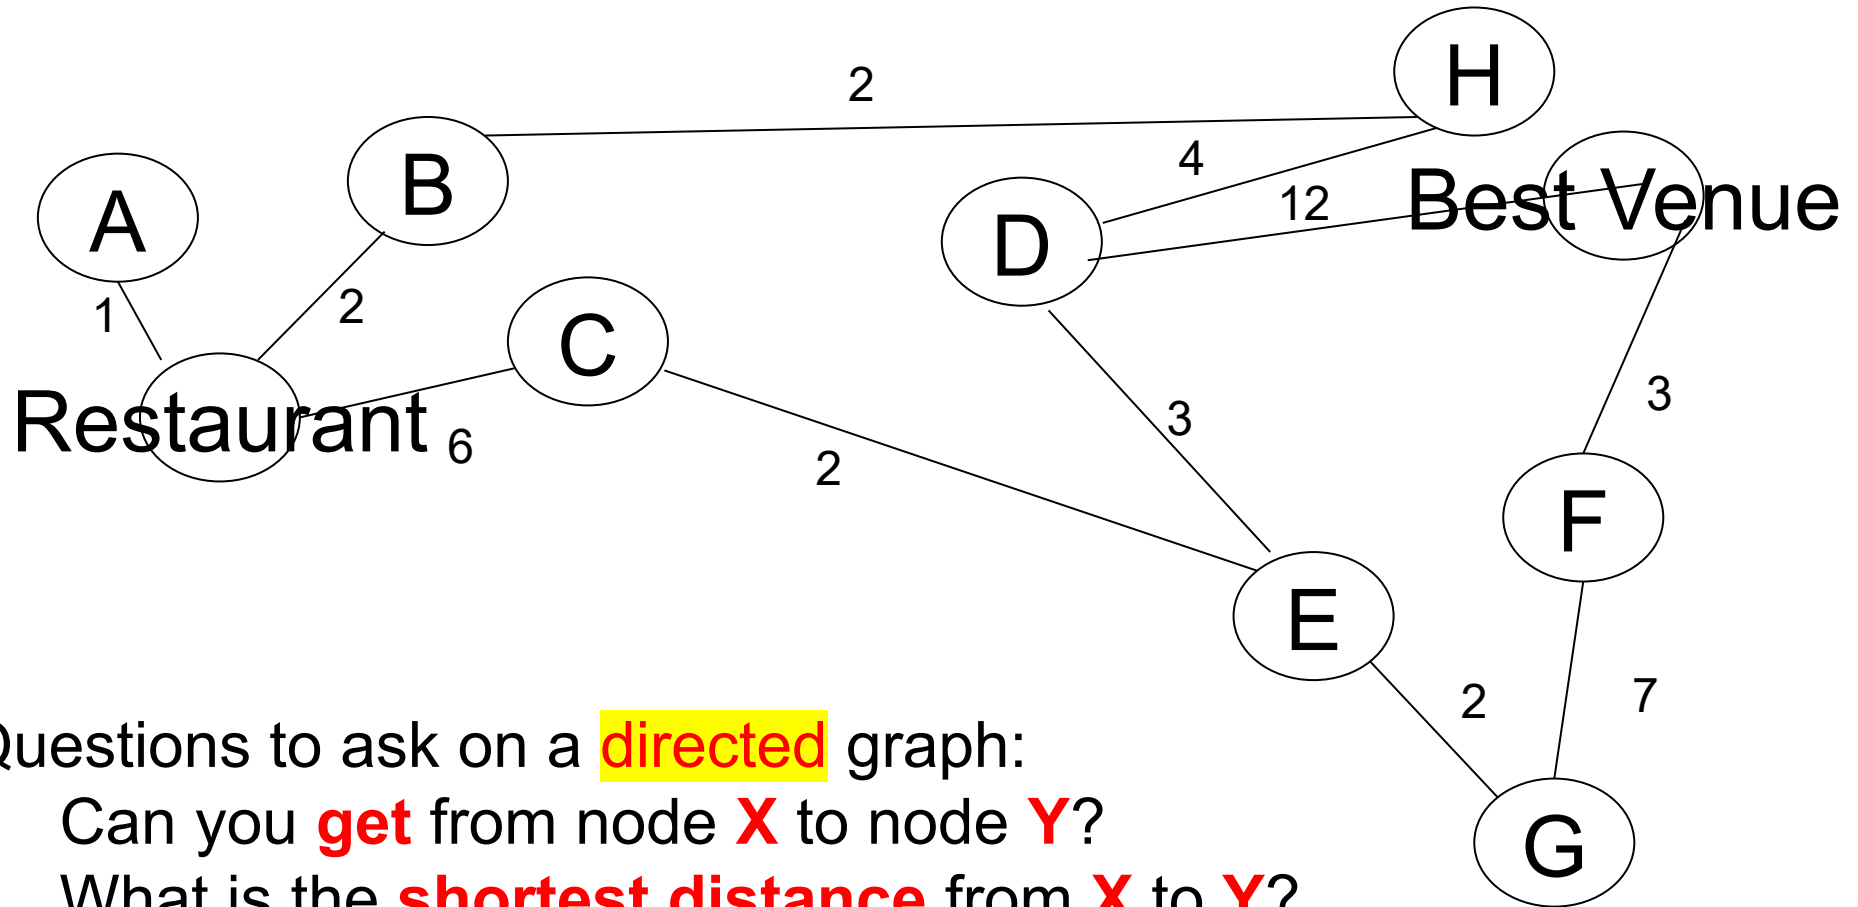


Questions to ask on a graph:

- Can you **get** from node **X** to node **Y**?
- What is the **shortest distance** from **X** to **Y**?
- What are the **shortest distances** from **X** to **any node**?
- What are the **shortest distances** from **any** node to **any** other node?



Example weighted graph



Questions to ask on a **directed** graph:

- Can you **get** from node **X** to node **Y**?
- What is the **shortest distance** from **X** to **Y**?
- What are the **shortest distances** from **X** to **any node**?
- What are the **shortest distances** from **any** node to **any** other node?



Single Source Shortest Path Problem

Given:

- Directed graph $G(V,E)$
- Source vertex s in V

Determine:

- **Shortest *distance*** path

from s to **every other** vertex in V



Brute force approach

For each vertex v_i :


- Enumerate all paths from s to v_i
- Calculate cost of each path $s \rightarrow \dots \rightarrow v_i$
- Pick minimum cost.

How many possible paths from s to v_i ?

- For a dense graph $O(V!)$
- $V=20$: 2432902008176640000 paths
- Not feasible!



Dijkstra's algorithm for single source shortest path




Greedy algorithm:

- Based on idea that **any subpath** along a shortest path **is also a shortest path**
- NodeA $\rightarrow \rightarrow \rightarrow \rightarrow$ NodeX \rightarrow NodeY
 - If **shortest path** from A to Y is **through X**,
 - then this path from A to X is also a **shortest path**

Dijkstra, E. W., *Numerische Mathematik* **1**: 269–271, 1959



Dijkstra's algorithm for single source shortest path



Greedy algorithm:

- Based on idea that **any subpath** along a shortest path **is also a shortest path**
- NodeA $\rightarrow \rightarrow \rightarrow \rightarrow$ NodeX \rightarrow NodeY
 - If **shortest path** from **A** to **Y** is through **X**,
 - then this path from **A** to **X** is also a **shortest path**

Assumes **no negative edges**, so:

$$\text{Distance}(A \rightarrow X) \leq \text{Distance}(A \rightarrow Y)$$




Dijkstra's algorithm for single source shortest path

Algorithm will give us a **shortest path tree**

Root = source node

- **Every node** is **connected** to the **root** through its **shortest path**

Image from R. Sedgewick, Lecture Notes
<http://www.cs.princeton.edu/courses/archive/fall05/cos226/lectures/shortest-path.pdf>





Dijkstra's Algorithm: Overview

For every vertex v and source s , maintain *estimate* **dist[v]** of *minimum distance* $\delta(s, v)$

dist[v]: length of a *known* path $s \rightarrow v$, but not necessarily the shortest path

- **dist[v]** $\geq \delta(s, v)$ Always
- When **dist[v]** $= \infty$, there is *no* estimate (*yet*)

Initially **dist[s]** $= 0$, all *other* **dist[v]** $= \infty$



Dijkstra's Algorithm: Overview

Process vertices **one-by-one**, **updating** **dist[v]** until **dist[v] = $\delta(s, v)$** , for every vertex v

- Along the way, **keep track of best path** information in array **pred[v]**

When algorithm finishes:

- Have shortest distances in **dist[]**
- Can reconstruct shortest path from **pred[]**



Relaxation: Updating estimated distances

Relaxation:

- **Estimate** the solution by answering an **easier problem** (relax the conditions)
- **dist[]** **Keeps updating** the relaxed estimate **until it is the solution** to the original problem


For shortest paths:

- **Estimate**: known distance of **best** path **so far**
- **Solution**: shortest possible distance



Relaxation: Updating estimated distances

Example:




`dist[v] : 35`
`pred[v] : n`
`pred[n] : s`

`dist[u] : 28`
`dist[v] : 30`
`pred[v] : u`



Relaxation: Updating estimated distances

Example:




```
void update_relax(u,v)
{
    if( dist[u] + edgeweight(u,v) < dist[v] )
    {
        dist[v] = dist[u] + edgeweight(u,v);
        pred[v] = u;
    }
}
```




Relaxation: Updating estimated distances

Example:



Note:

```
pred[v] = u;
```

```
pred[u] = n;
```

```
...
```

```
pred[j] = s;
```

Reconstruct path $s \rightarrow v$ going **backwards** through **pred[]**



Dijkstra's algorithm: successive relaxations

How do we **pick the next node** to look at?

Process vertices in order of **estimated closeness** to source, value of **`dist[v]`**

Priority queue to store vertex v and **`dist[v]`** value



Dijkstra's algorithm (C-ish pseudocode)

```
/* Find shortest paths in graph G from source s*/  
/* vertices identified by number for convenience */  
  
void dijkstra(int** G, int s)  
{  
    int dist[Vsize], pred[Vsize];  
    initialize(G, s, pred, dist);  
    run(G, s, pred, dist);  
  
    reconstruct(s, pred, dist);  
}
```



Dijkstra's algorithm (C-ish pseudocode)

```
void initialize(int** G, int Vsize, int s, int* pred, int* dist)
{
    int i;
    for( i = 0; i < Vsize; i++)
        dist[i] = MAX_INT;
    dist[s] = 0;
    for( i = 0; i < V; i++)
        pred[i] = NULL;
}
```



Dijkstra's algorithm(C-ish pseudocode)

```
void run(int** G, int Vsize, int s, int* pred, int* dist)
{
    pq_node_t* pq;
    int u, v;
    pq = makePQ(G); /* vertices into min PQ, dist as priority */

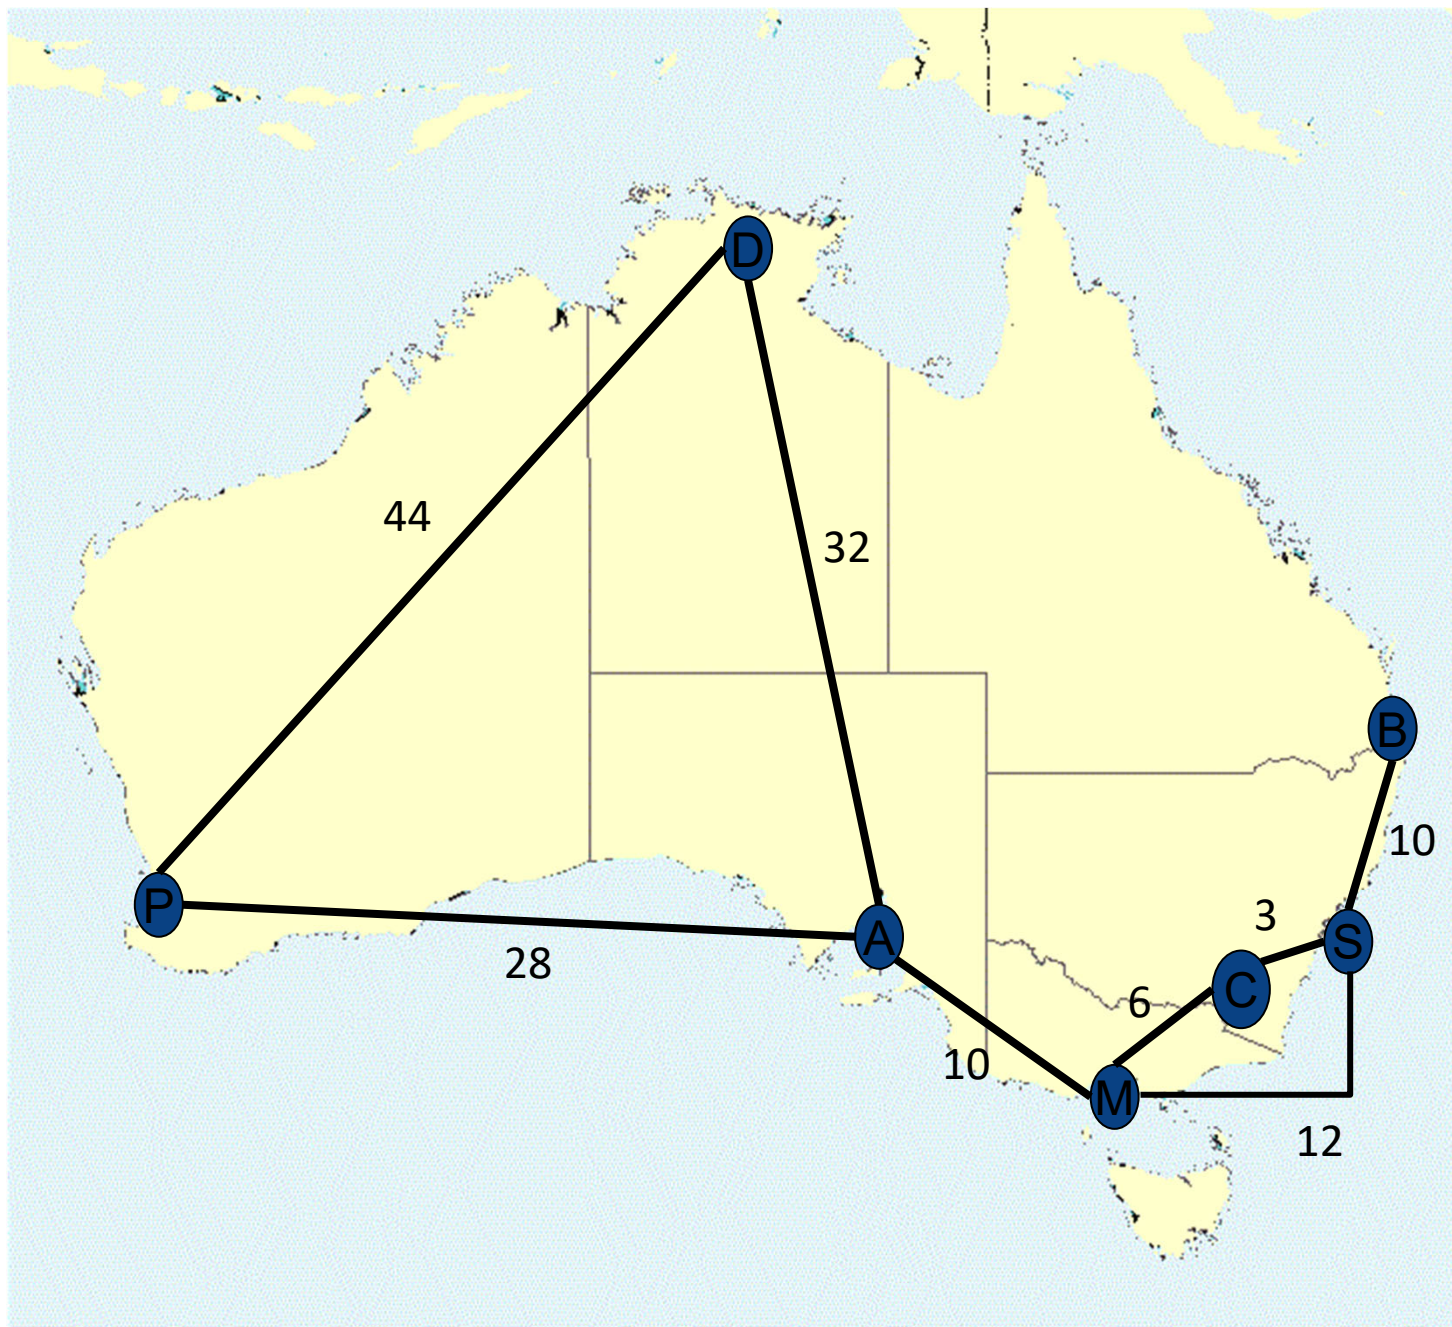
    while( !emptyPQ(pq) )
    {
        u = deletemin(pq);
        /* At this point vertex u has been processed,
        * i.e. dist[u] = delta(s,u) = shortest path to u found */

        for(/*each v conneted to u */)
            if(dist[u] + edgeweight(u,v) < dist[v])
                update(v, pred, dist, pq);
    }
}
```



Dijkstra's algorithm (C-ish pseudocode)

```
void update(int v, int* pred, int* dist, pq_node_t* pq)
{
    dist[v] = dist[u] + edgeweight(u,v);
    pred[v] = u;
    decreaseweight(pq, v, dist[v]);
}
```





Dijkstra's algorithm (Example)

M C S B A P D

D =


Pred =

PQ =

u=

v ∈ { }


```
if(dist[u] + edgeweight(u,v) < dist[v])  
    update(v, pred, dist, pq);
```





Dijkstra's algorithm: example

Example:




Step 0: Initialization.

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞
$pred[v]$	nil	nil	nil	nil	nil
$color[v]$	W	W	W	W	W

Priority Queue:

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞




Step 1: As $Adj[s] = \{a, b\}$, work on a and b and update information.

v	s	a	b	c	d
$d[v]$	0	2	7	∞	∞
$pred[v]$	nil	s	s	nil	nil
$color[v]$	B	W	W	W	W

Priority Queue:

v	a	b	c	d
$d[v]$	2	7	∞	∞

From lecture notes by Mordechai Golin, Univ Science and Technology, Hong Kong:
<http://www.cse.ust.hk/faculty/golin/COMP271Sp03/Notes/MyL09.pdf>




Step 2: After Step 1, a has the minimum key in the priority queue. As $Adj[a] = \{b, c, d\}$, work on b , c , and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	10	7
$pred[v]$	nil	s	a	a	a
$color[v]$	B	B	W	W	W

Priority Queue:

v	b	c	d
$d[v]$	5	10	7




Step 3: After Step 2, b has the minimum key in the priority queue. As $Adj[b] = \{a, c\}$, work on a , c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	W	W

Priority Queue:

v	c	d
$d[v]$	6	7




Step 4: After Step 3, c has the minimum key in the priority queue. As $Adj[c] = \{d\}$, work on d and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	W

Priority Queue:

v	d
$d[v]$	7



Step 5: After Step 4, d has the minimum key in the priority queue. As $Adj[d] = \{c\}$, work on c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	B

Priority Queue: $Q = \emptyset$.



Dijkstra's algorithm(ANALYSIS)

```
void run(int** G, int Vsize, int s, int* pred, int* dist)
{
    pq_node_t* pq; /****** Assuming PQ is a minheap */
    int u, v;
    pq = makePQ(G); /****** big-O()???? */

    while( !emptyPQ(pq) )
    {
        u = deletemin(pq); /****** big-O()???? */
        for(/*each v conneted to u */)
            if(dist[u] + edgeweight(u,v) < dist[v])
                update(v, pred, dist, pq); /****** big-O()???? */
    }
}
```



Dijkstra's Algorithm: Analysis

Cost depends on implementation of PQ

Using a heap:

- **makePQ** () $O(V)$
- V * **deletemin** () operations @ $O(\log V)$
- $O(E)$ **decreaseweight** () ops @ $O(\log V)$
- Total: $O((V+E) \log V)$



Dijkstra's Algorithm: Limitations

Assumes **no negative edges**:

- Good for physical distances
- Distances are static

Negative edges:

- Use **Bellman-Ford algorithm**
- **Cannot** deal with **negative cycles**
- $O(V * E)$



Dijkstra's Algorithm: Limitations

Negative *cycles*:

- What is the **shortest** path?
- Problem is not well-formed, intractable
- Bellman-Ford detects negative cycles (algorithm does terminate, stops keeps shortening paths)

Tutorial:

<https://www.dyclassroom.com/graph/detecting-negative-cycle-using-bellman-ford-algorithm>



Applications

More applications

- Robot navigation
- Texture mapping
- Typesetting in TeX
- Urban traffic planning
- Optimal pipelining of VLSI chip
- Telemarketer operator scheduling
- Routing of telecommunications messages.
- Network routing protocols (OSPF, BGP, RIP)
- Exploiting arbitrage opportunities in currency exchange
- Optimal truck routing through given traffic congestion pattern



Negative Cycle Detection: Arbitrage

Common example in CS materials is arbitrage:

- currency 1 \rightarrow currency 2 \rightarrow currency 3 \rightarrow currency 1'
- If currency 1' > currency 1, you have made money
- Model problem as a graph:
 - Vertices = currency
 - Edges = $-\log_2(\text{exchange rate})$
 - Detect negative cycle and change money \rightarrow get rich!

Not realistic!

- D.J.Fenn *et al.*, “The Mirage of Triangular Arbitrage in the Foreign Currency Exchange Market”, *Int. J. Theoretical and Applied Finance* **12**(8), 1105-1123, 2009.



Edsger W. Dijkstra

- The question of whether **computers can think** is like the question of whether **submarines can swim**
- *Computer science is no more about **computers** than **astronomy** is about **telescopes***
- How do we convince people that in programming **simplicity** and **clarity** —in short: what mathematicians call "**elegance**"— are not a dispensable luxury, but a crucial matter that decides between success and failure?

Elegance is not a dispensable luxury but a quality that decides between success and failure

Turing award 1972



A very nice explanation of Dijkstra's algorithm by Mordechai Golin can be found at

<http://www.cse.ust.hk/faculty/golin/COMP271Sp03/Notes/MyL09.pdf>