

Data Structures and Algorithms

11. Heuristic Search Algorithms

Basic Stuff You're Gonna Need to Search over a Graph
Where To Search Next?

Nir Lipovetzky



THE UNIVERSITY OF
MELBOURNE

Path-finding in graphs

Search algorithms over directed graphs:

- The **search nodes** (vertex) of the graph represent some information
- The edges (v, v') capture transitions

One way to understand search algorithms is to think about **path-finding** over **graphs**.

Classification of Search Algorithms

Blind search vs. heuristic (or informed) search:

- **Blind search algorithms:** Only use the basic ingredients for general search algorithms.
 - e.g., *Depth First Search (DFS)*, *Breadth-first search (BrFS)*, *Uniform Cost (Dijkstra)*, *Iterative Deepening (ID)*
- **Heuristic search algorithms:** Additionally use **heuristic functions** which estimate the distance (or remaining cost) to reach an target vertex.
 - e.g., *A**, *IDA**, *Hill Climbing*, *Best First*, *WA**, *DFS B&B*, *LRTA**, ...

Systematic search vs. local search:

- **Systematic search algorithms:** Consider a large number of search nodes simultaneously.
- **Local search algorithms:** Work with one (or a few) candidate solutions (search nodes) at a time.
 - This is not a black-and-white distinction; there are *crossbreeds* (e.g., **enforced hill-climbing**).

Before We Begin

Blind search vs. informed search:

- **Blind search** does not require any input beyond the graph.
 - **Pros and Cons?** Pro: No additional work for the programmer. Con: It's not called "blind" for nothing ... same expansion order regardless what the problem actually is.
- **Informed search** requires as additional input a **heuristic function h** that maps nodes to estimates of their **distance**.
 - **Pros and Cons?** Pro: Typically more effective in practice. Con: Somebody's gotta come up with/implement h .

Before We Begin, ctd.

Blind search strategies we'll discuss:

- **Breadth-first search**. Advantage: time complexity.
Variant: **Uniform cost search**.
- **Depth-first search**. Advantage: space complexity.
- **Iterative deepening search**. Combines advantages of breadth-first search and depth-first search. Uses **depth-limited search** as a sub-procedure.

Heuristic Search Algorithms: Systematic

→ Heuristic search algorithms are the most common and overall most successful algorithms for path-finding in graphs.

Systematic heuristic search algorithms:

- Greedy best-first search.
- Weighted A^* .
- A^* .
- IDA^* , depth-first branch-and-bound search, ...


Heuristic Search Algorithms: Local

→ Heuristic search algorithms are the most common and overall most successful algorithms for path-finding.

Local heuristic search algorithms:

- [Hill-climbing](#).
- Beam search, tabu search, genetic algorithms, simulated annealing, ...

Heuristic Search: Basic Idea



→ Heuristic function h estimates the cost of an optimal path to the goal; search gives a preference to explore states with small h .

Heuristic Functions

Heuristic searches require a heuristic function to estimate remaining cost:

Definition (Heuristic Function). A **heuristic function**, short **heuristic**, is a function $h : S \mapsto \mathbb{R}_0^+ \cup \{\infty\}$. Its value $h(s)$ for a state s is referred to as the state's **heuristic value**, or **h -value**.

Definition (Remaining Cost, h^*). For a state $s \in S$, the state's **remaining cost** is the cost of a shortest path from s to a goal state g , or ∞ if there exists no path for s . The **perfect heuristic**, written h^* , assigns every $s \in S$ its remaining cost as the heuristic value.

→ Recall that a state s is equivalent to a vertex v in the graph!

Heuristic Functions: Discussion

What does it mean to “estimate remaining cost”?

- For many heuristic search algorithms, h does not need to have any properties for the algorithm to “work” (= be correct and complete).
→ h is *any* function from states to numbers . . .
- Search **performance** depends crucially on “how well h reflects h^* ”!!
→ This is informally called the **informedness** or **quality** of h .
- For some search algorithms, like A^* , we can *prove* relationships between formal quality properties of h and search efficiency (mainly the number of expanded nodes).
- For other search algorithms, “it works well in practice” is often as good an analysis as one gets.

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Any other property of h that search performance crucially depends on?

“... and on the computational overhead of computing h !!”

Extreme cases:

- $h = h^*$: Perfectly informed; computing it = solving the original problem in the first place.
- $h = 0$: No information at all; can be “computed” in constant time.

→ Successful heuristic search requires a good trade-off between h 's informedness and the computational overhead of computing it.

→ **This really is what research is all about!** Devise methods that yield good estimates at reasonable computational costs.

Properties of Heuristic Functions

Definition. The heuristic is called:

- **Admissible** if $h(s) \leq h^*(s)$ for all $s \in S$;

Other properties not discussed in this course: [Goal-Aware](#), [Consistent](#) and [Safe](#).

Motivation: For You

Imagine ...

You have completed your studies, and have been hired by some company in the software-making business.

Boss: *Hey you, here's that problem. Solve it.*

You (thinking): *Hm, I think heuristic search might work.*

You (thinking): *Hm, I need a heuristic function. How?!?*

→ “Relax”ing is a methodology to construct heuristic functions.

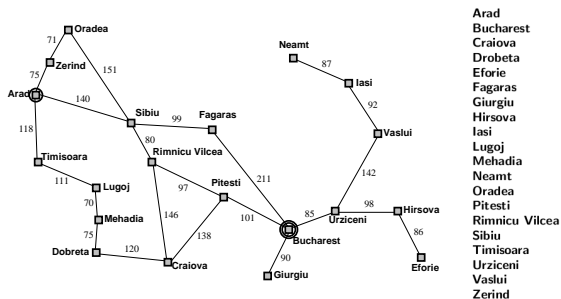
How to Relax Informally

How To Relax:

- You have a problem, \mathcal{P} , whose perfect heuristic h^* you wish to estimate.
- You define a **simpler problem**, \mathcal{P}' , whose perfect heuristic h'^* can be used to **estimate h^*** .
- You define a transformation, r , that **simplifies** instances from \mathcal{P} into instances \mathcal{P}' .
- Given $\Pi \in \mathcal{P}$, you estimate $h^*(\Pi)$ by $h'^*(r(\Pi))$.

→ Relaxation means to simplify the problem, and take the solution to the simpler problem as the heuristic estimate for the solution to the actual problem.

Relaxation in Route-Finding



How to derive straight-line distance by relaxation?

- Problem \mathcal{P} : Route finding.
- Simpler problem \mathcal{P}' : Route finding for birds.
- Perfect heuristic h'^* for \mathcal{P}' : Straight-line distance.
- Transformation r : Pretend you're a bird.

Greedy Best-First Search

Greedy Best-First Search (with duplicate detection)

```
open := new priority queue ordered by ascending  $h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min()}$  /* get best state */
    if  $\text{state}(\sigma) \notin \text{closed}$ : /* check duplicates */
         $\text{closed} := \text{closed} \cup \{\text{state}(\sigma)\}$  /* close state */
        if is-goal( $\text{state}(\sigma)$ ): return extract-solution( $\sigma$ )
        for each ( $s' \in \text{succ}(\text{state}(\sigma))$ ): /* expand state */
             $\sigma' := \text{make-node}(\sigma, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ ): return unsolvable
```


Greedy Best-First Search: Remarks

Properties:

- **Complete?** Yes (Due to duplicate detection.)
- **Optimal?** No.¹

Implementation:

- Priority queue: e.g., a [min heap](#).
- “Check Duplicates”: Could already do in “expand state”; done here after “get best state” *only* to more clearly point out relation to A^* .

¹ Even for perfect heuristics! E.g., say the start state has two transitions to goal states, one of which costs a million bucks while the other one is free. Nothing keeps Greedy Best-First Search from choosing the bad one.

A*


A* (with duplicate detection and re-opening)

```

open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$  /* maps states to numbers */
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        /* re-open if better g; note that all  $\sigma'$  with same state but worse g
           are behind  $\sigma$  in open, and will be skipped when their turn comes */
        closed := closed  $\cup$  {state( $\sigma$ )}
        best-g(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each ( $s'$ )  $\in$  succ(state( $\sigma$ )):
             $\sigma' := \text{make-node}(\sigma, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```

A*: Example



A*: Terminology

- **f -value** of a state: defined by $f(s) := g(s) + h(s)$.
- **Generated nodes**: Nodes inserted into *open* at some point.
- **Expanded nodes**: Nodes σ popped from *open* for which the test against *closed* and *distance* succeeds.
- **Re-expanded nodes**: Expanded nodes for which $state(\sigma) \in closed$ upon expansion (also called **re-opened** nodes).

A*: Remarks

Properties:

- **Complete?** Yes (Even without duplicate detection.)
- **Optimal?** Yes, for admissible heuristics. (Even without duplicate detection.)

Implementation:

- Popular method: break ties ($f(s) = f(s')$) by smaller h -value.
- Common, hard to spot bug: check duplicates at the wrong point.
- **Our implementation is optimized for readability not for efficiency!**

Hill-Climbing

Hill-Climbing

```

 $\sigma$  := make-root-node(init())
forever:
  if is-goal(state( $\sigma$ )):
    return extract-solution( $\sigma$ )
   $\Sigma' := \{ \text{make-node}(\sigma, s') \mid (s') \in \text{succ}(\text{state}(\sigma)) \}$ 
   $\sigma :=$  an element of  $\Sigma'$  minimizing  $h$  /* (random tie breaking) */

```

Remarks:

- Is this complete or optimal? No.
- Can easily get stuck in **local minima** where immediate improvements of $h(\sigma)$ are not possible.
- Many variations: tie-breaking strategies, restarts, ...

Properties of Search Algorithms

ID: Iterative Deepening (not included in this lecture)

	DFS	BrFS	ID	A*	HC
Complete	No	Yes	Yes	Yes	No
Optimal	No	Yes*	Yes	Yes	No
Time	∞	b^d	b^d	b^d	∞
Space	$b \cdot d$	b^d	$b \cdot d$	b^d	b

- Parameters: d is solution depth; b is branching factor
- Breadth First Search (BrFS) optimal when costs are uniform
- A* optimal when h is **admissible**; $h \leq h^*$
- Search Algorithms visualization:
<https://www.youtube.com/watch?v=rZHtHJlJa2w>
- More Search Algorithms visualization: <http://www.redblobgames.com/pathfinding/a-star/introduction.html>
- Interactive Grid Path Finding visualization:
<https://qiao.github.io/PathFinding.js/visual/>

Applications of Path finding

Many problems can be **understood and solved** as path-finding search problems over graphs!

- Google self driving car

<https://youtu.be/qXZt-B7iUyw>

- Computer Games

<https://www.youtube.com/watch?v=DlkMs4ZHHr8>

- Base algorithm (A^*) running “Best AI in Computer Games”

<https://youtu.be/10Xb8mg9IVw?t=7m45s>

- Local search, Evolution

<https://www.youtube.com/watch?v=bBt0imn77Zg>

Questionnaire, ctd.

Question!

Is informed search always better than blind search?

(A): Yes.

(B): No.

→ In greedy best-first search, the heuristic may yield larger search spaces than uniform-cost search. E.g., in path finding, say you want to go from Melbourne to Sydney, but $h(\text{Perth}) < h(\text{Canberra})$.

→ In A^* with an admissible heuristic and duplicate checking, we cannot do worse than uniform-cost search: $h(s) > 0$ can only reduce the number of states we must consider to prove optimality.

→ Also, in the above example, A^* doesn't expand Perth with *any* admissible heuristic, because $g(\text{Perth}) > g(\text{Sydney})$!

→ "Trusting the heuristic" has its dangers! Sometimes g helps to reduce search.

Summary

Search algorithms mainly differ in **order of node expansion**:

- **Blind** vs. **heuristic** (or **informed**) search.
- **Systematic** vs. **local** search.

Summary (ctd.)

- **Search strategies** differ (amongst others) in the order in which they **expand search nodes**, and in the way they use **duplicate elimination**. Criteria for evaluating them are **completeness**, **optimality**, **time complexity**, and **space complexity**.
- **Breadth-first search** is optimal but uses exponential space; **depth-first search** uses linear space but is not optimal. **Iterative deepening search** combines the virtues of both.

Summary (ctd.)

Heuristic Functions: Estimators for remaining cost.

- Usually: The more informed, the better performance.
- Desiderata: admissible, informed
- The ideal: Perfect heuristic h^* .

Heuristic Search Algorithms:

- Most common algorithms if you don't care about optimality:
 - Greedy best-first search.
 - Weighted A^* .
 - hill-climbing.
- Most common algorithm if you care about optimality:
 - A^* .
 - IDA^* .

Search Algorithms for Path Finding in Directed Graphs

■ Blind search/Brute force algorithms

- Goal plays **passive** role in the search
- *e.g., Depth First Search (DFS), Breadth-first search (BrFS), Uniform Cost (Dijkstra), Iterative Deepening (ID)*

■ Informed/Heuristic Search Algorithms

- Goals plays **active** role in the search through **heuristic function** $h(s)$ that estimates cost from s to the goal
- *e.g., A^* , IDA^* , Hill Climbing, Best First, WA^* , DFS B&B, $LRTA^*$, ...*